Oregon State University

Apple Orchards: A multitask **multiagent** approach

Alejo Velasquez, Nigel Swenson and Keegan Nave

ROB538 Fall'22 - Project Presentation

Credits: https://www.therobotreport.com/abundant-robotics-shuts-down-fruit-harvesting-business/

Motivation

Apple orchards are promising for robots.

Orchards have multiple tasks that need to be performed.

How to plan for multiple agents that handle different **sequential** tasks?

Credits: https://www.abc.net.au/news/2021-05-03/robots-almost-ready-to-help-pick-fruit/100107796

Source: https://usapple.org/news-resources/apple-stages-from-the-tree-to-the-grocery-store

Background

Hybrid (human-robot) apple picking

Different skills, centralized

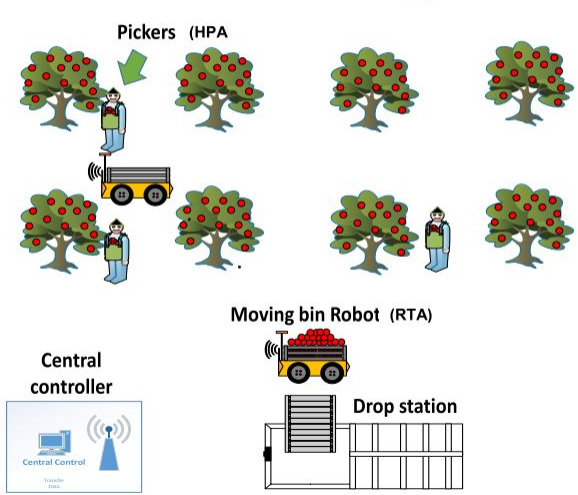

(Salah, 2019)

Vineyard irrigation

Same skill, many agents

(Thayer et al., 2020)

Problem Definition

Problem Space

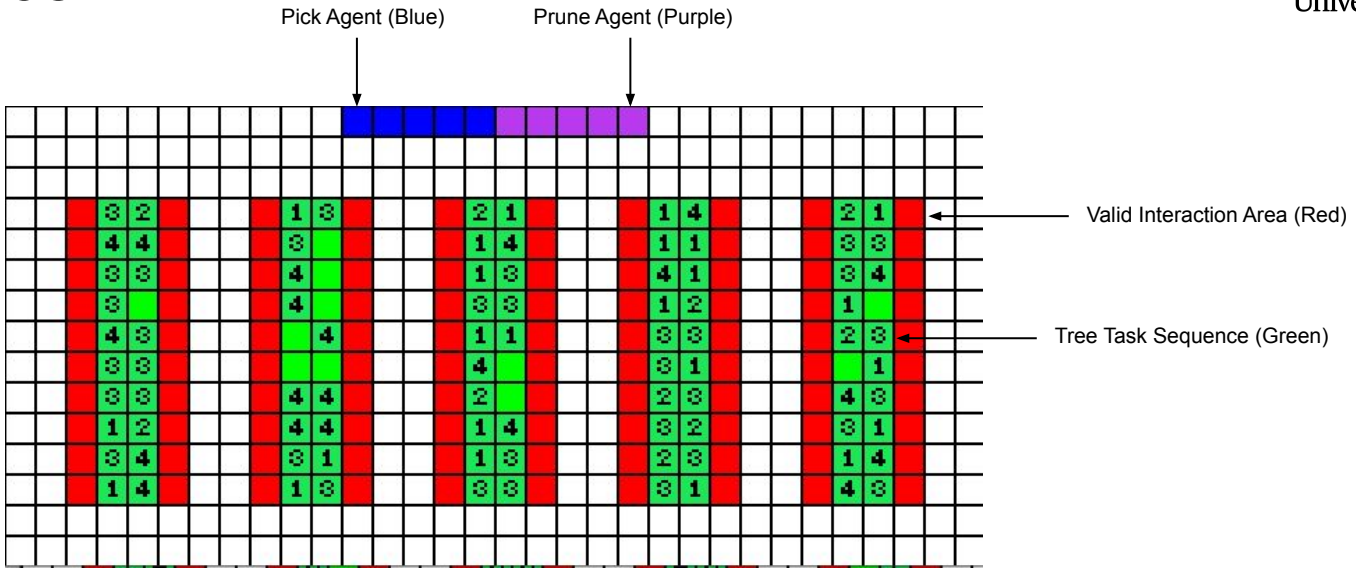

$$TaskSequences = \begin{cases} 1 : \text{Pick} \rightarrow \text{Complete} \\ 2 : \text{Prune} \rightarrow \text{Complete} \\ 3 : \text{Pick} \rightarrow \text{Prune} \rightarrow \text{Complete} \\ 4 : \text{Prune} \rightarrow \text{Pick} \rightarrow \text{Complete} \\ None : \text{Complete} \end{cases}$$

$$S_i = \{Agent_i(x, y), ValidTree_1, ValidTree_2, ..., ValidTree_n\}$$

$$A = \{up, down, left, right, interact\}$$

Goals

- Investigate tightly coupled agents in Agricultural settings.
- Improve performance on sequential task spaces.
- Maximize the amount of pick and prune tasks that are completed.

Reward Structure

Local Reward

$$L = \left\{ 10 \text{ when interacting}, -1 \text{ otherwise} \right\}$$

Global Reward

$$G = \frac{TreesPruned}{TotalLeaves} \times \frac{ApplesPicked}{ApplestoPick}$$

Potential Local Reward

Difference Reward 1: Random Action

Difference Reward 2: Random Action with 1 step look ahead

Approach 1

Q learning - what **didn't** work

Oregon State University

Let's start simple

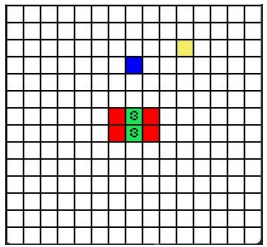

2 rows, 1 col, 2 agents
1k steps, gamma 0.9
alpha 0.05, epsilon*0.99

Local Rewards

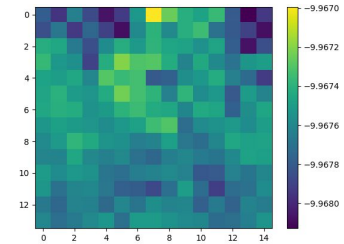

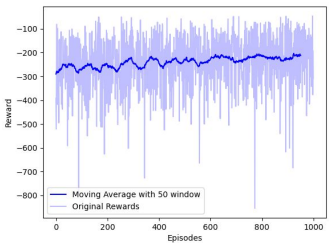

Global Rewards

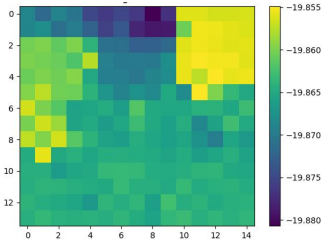

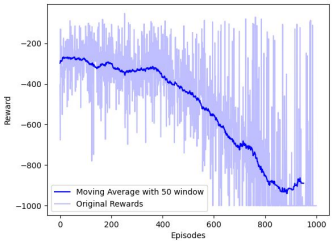

Diff Rewards 1

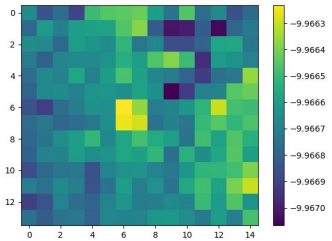

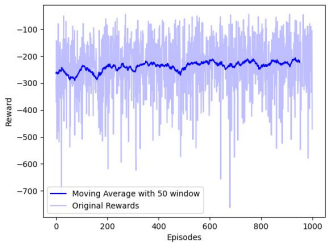

Q learning - what **did** work

Sequential tasks can be thought as **tightly-coupled**.

What if a tree is **not** useful for **me**, but it is for **others**?

```
i: current agent
j: other agents

...
if i.reward < 0:
   for j in agents:
      # Let each agent assume the current agent's position
      j.valid_moves = get_valid_moves(i.cur_pose, j.type)

      if "interact" in j.valid_moves:
         total_rewards += 10
...
```

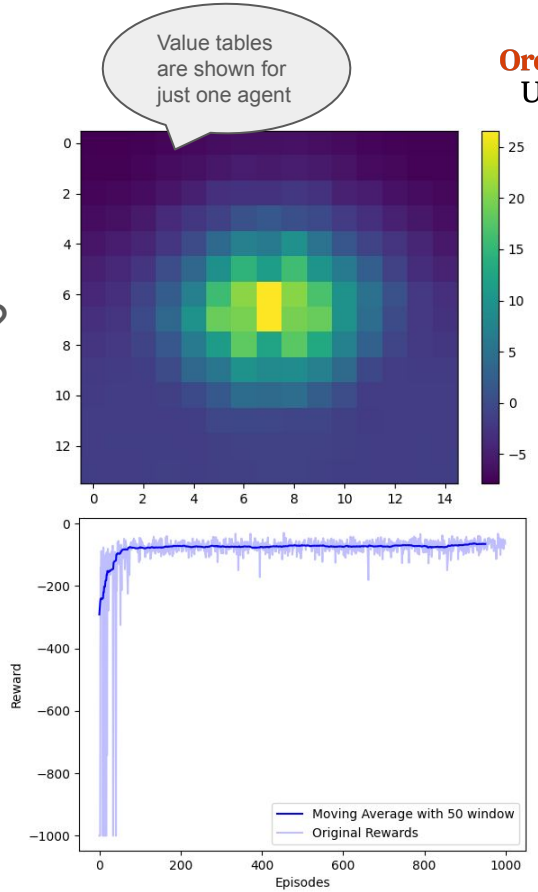

Q tables - up to a certain point

6 agents, **2 columns** and 2 rows

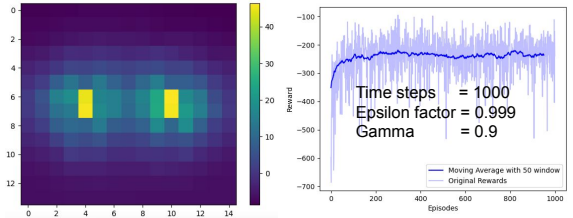

Sequence (3: Prune → 2: Pick → 0: Done)

6 agents, **5 columns** and 2 rows

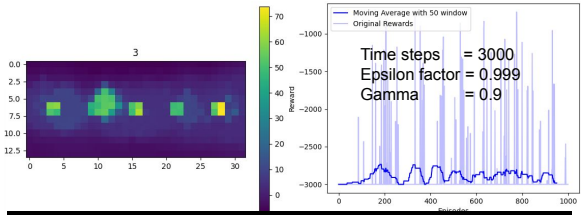

Sequence (3: Prune → 2: Pick → 0: Done)

Sensitive to number of agents and size of the orchard

Approach 2

Soft Actor Critic

Replace Q table with Actor and Critic Networks

Can handle larger state spaces - don't have to memorize every point

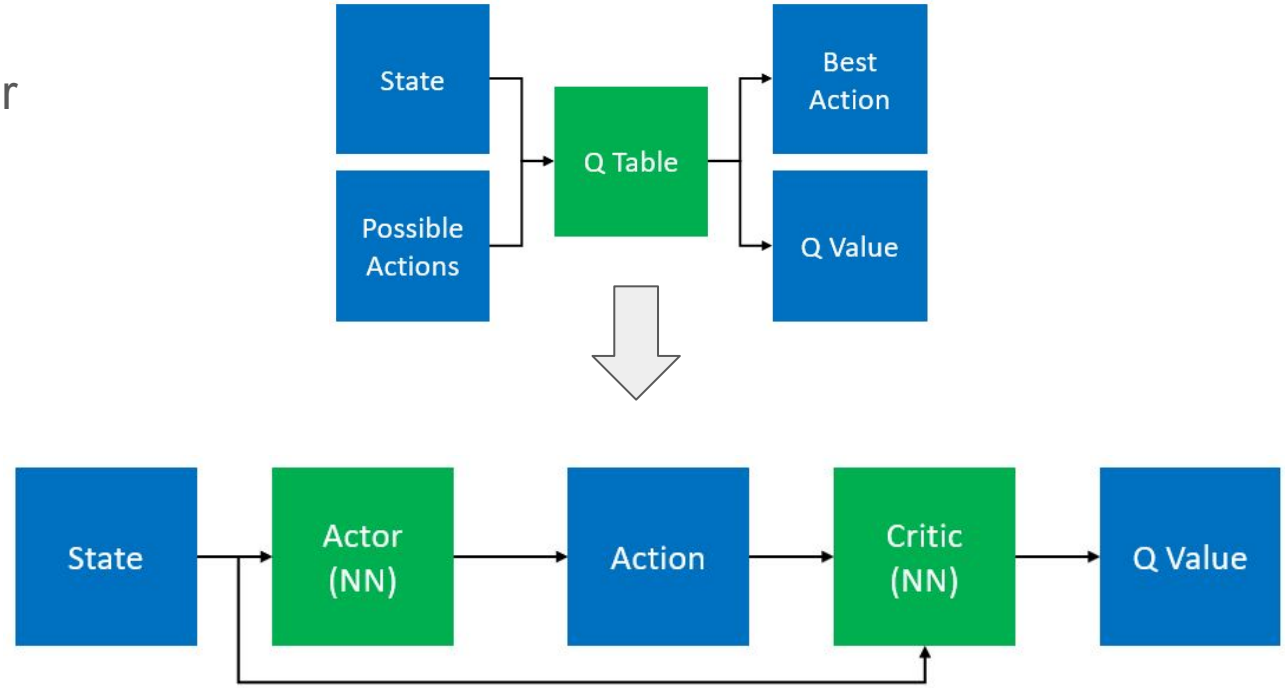

Results SAC Same Seed

1 pick agent

1 prune agent

12x9 orchard

100 steps/episode

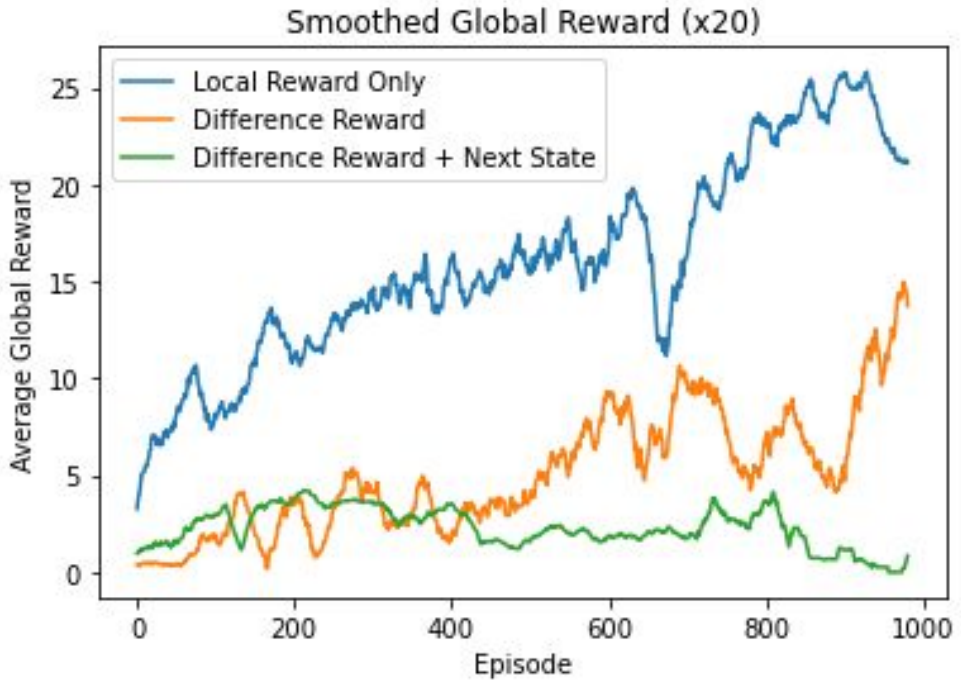

Results SAC Different Seed

1 pick agent

1 prune agent

12x9 orchard

100 steps/episode

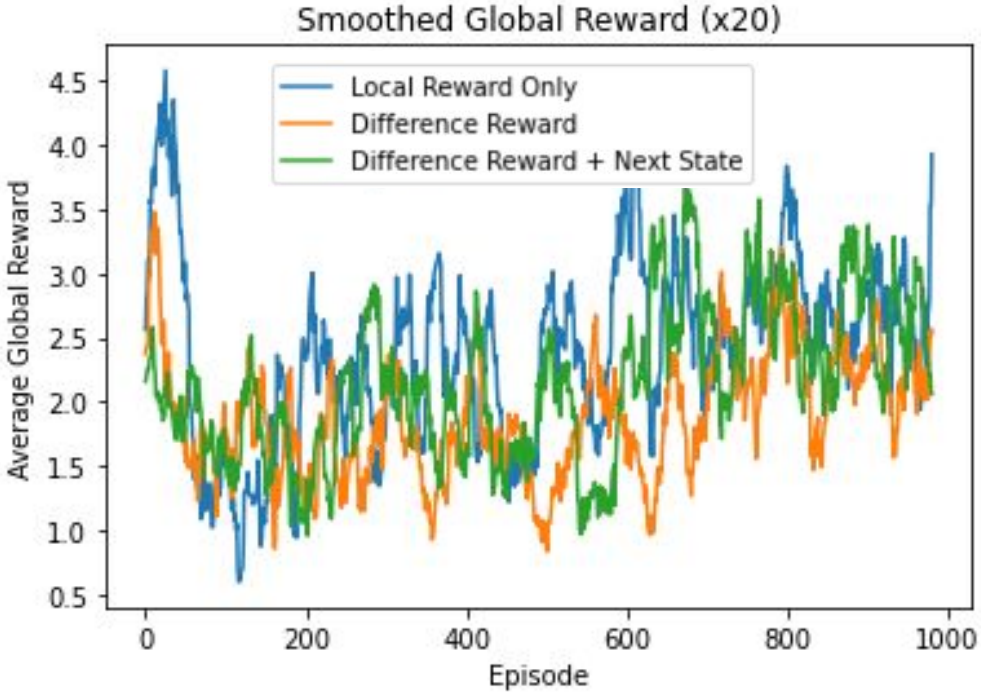

Conclusions

Conclusions

- Qtables were useful for smaller grids. **With greater grids, specially with more columns, the system didn’t converge.**

- Local rewards outperform difference reward with current global baseline
  - Worth changing our global system to add the fractions to encourage individual actions

Future Work

Oregon State University

Short term (paperwise)

- Explore other **reward shaping** techniques in Q-table approach.
- Explore **higher number** of sequential tasks.
- Compare with other **benchmarks** (e.g. team of pruners go first, then team of pickers)

Long term (research)

- Explore impact of **orchard structure.**
- Explore impact of **time** between tasks.

Thank you

Questions?